

Generalist (G)


all subtasks · domain knowledge



*one model, all three
label sets; the prompt
selects the subtask*

Specialist (S)


one subtask · full label set



*one model per sub-
task; narrower sig-
nal, different errors*

Minority-Class Spec. (MCS)

one subtask · rare labels only



*drop the most frequent labels
until the dropped set passes 50 %*

the rule per subtask

ST1 drop `physical_goods` 47 % + `digital_content_or_services` 46 % = **93 %**, keep 3 of 5 labels

ST2 drop `apps` 29 % + `hardware_electronics` 22 % = **51 %**, keep 10 of 12 labels

ST3 drop `misleading_claim` **54 %**, keep 7 of 8 labels

ST1 is single-label, so dropping a label drops its instances; ST2 and ST3 are multi-label, so only the label columns go and the instances stay.